

ISTM 6212 - Week 5

RDBMS schema design, joins, etc.

Daniel Chudnov, dchud@gwu.edu

Agenda

- ❖ Project 01 - Reviews
- ❖ RDBMS DML: SQL JOINS, subqueries
- ❖ RDBMS DDL: schema, normal forms, E-R models
- ❖ RDBMS DML: SQL CREATE, INSERT, UPDATE, DELETE
- ❖ RDBMS in practice: transactions, functions, triggers
- ❖ Exercise 03

Project 01 - Reviews

RDBMS DML - SQL JOINs


Symbol (Name)	Example of Use
σ (Selection)	$\sigma \text{ salary} \geq 85000 \text{ (instructor)}$
	Return rows of the input relation that satisfy the predicate.
Π (Projection)	$\Pi_{ID, salary} \text{ (instructor)}$
	Output specified attributes from all rows of the input relation. Remove duplicate tuples from the output.
\bowtie (Natural Join)	$\text{instructor} \bowtie \text{department}$
	Output pairs of rows from the two input relations that have the same value on all attributes that have the same name.
\times (Cartesian Product)	$\text{instructor} \times \text{department}$
	Output all pairs of rows from the two input relations (regardless of whether or not they have the same values on common attributes)
\cup (Union)	$\Pi_{name}(\text{instructor}) \cup \Pi_{name}(\text{student})$
	Output the union of tuples from the two input relations.

SQL JOIN


- ❖ Cartesian product by default
- ❖ further specified by common attributes with ON
- ❖ multiple tables, multiple pairs of common attributes
- ❖ still all based on set operations

SQL JOIN types


❖ INNER JOIN




❖ LEFT OUTER JOIN



❖ RIGHT OUTER JOIN



❖ FULL OUTER JOIN




```

SELECT DISTINCT
    display_call_no,
    item_status_desc,
    item_status.item_status,
    perm_location.location_display_name as PermLocation,
    temp_location.location_display_name as TempLocation,
    mfhd_item.item_enum,
    mfhd_item.chron,
    item.item_id,
    item_status_date,
    to_char(CIRC_TRANSACTIONS.CHARGE_DUE_DATE, 'yyyy-mm-dd') AS DUE,
    library.library_display_name,
    holding_location.location_display_name as HoldingLocation
FROM bib_master
JOIN library ON library.library_id = bib_master.library_id
JOIN bib_mfhd ON bib_master.bib_id = bib_mfhd.bib_id
JOIN mfhd_master ON mfhd_master.mfhd_id = bib_mfhd.mfhd_id
JOIN library ON bib_master.library_id = library.library_id
JOIN location holding_location
    ON mfhd_master.location_id = holding_location.location_id
LEFT OUTER JOIN mfhd_item
    ON mfhd_item.mfhd_id = mfhd_master.mfhd_id
LEFT OUTER JOIN item
    ON item.item_id = mfhd_item.item_id
LEFT OUTER JOIN item_status
    ON item_status.item_id = item.item_id
LEFT OUTER JOIN item_status_type
    ON item_status.item_status = item_status_type.item_status_type
LEFT OUTER JOIN location perm_location
    ON perm_location.location_id = item.perm_location
LEFT OUTER JOIN location temp_location
    ON temp_location.location_id = item.temp_location
LEFT OUTER JOIN circ_transactions
    ON item.item_id = circ_transactions.item_id
WHERE bib_master.bib_id = %s
AND mfhd_master.suppress_in_opac != 'Y'
ORDER BY PermLocation, TempLocation, item_status_date desc

```


- ❖ if your work requires you to do this often, you will become good at it quickly.
- ❖ just think in sets!

RDBMS DML - subqueries

SQL subqueries

- ❖ remember: tables are relations and query results are relations
- ❖ relations are sets of attribute value instances
- ❖ we can use one relation to specify attribute conditions for another

SQL subquery example

- ❖ Query 1: SELECT DISTINCT ident FROM person;
- ❖ Query 2: SELECT * FROM survey;
- ❖ use Query 1 result within Query 2:

```
SELECT * FROM survey  
WHERE person IN  
(SELECT DISTINCT ident FROM person);
```



```

def get_related_bibids_by_isbn(item):
    if 'isbn' not in item or len(item['isbn']) == 0:
        return []

    binds = ','.join(['%s'] * len(item['isbn']))

    q = '''
SELECT DISTINCT bib_index.bib_id, bib_text.title
FROM bib_index, bib_master, bib_text
WHERE bib_index.bib_id=bib_master.bib_id
AND bib_master.suppress_in_opac='N'
AND bib_index.index_code IN ('020N','020A','ISB3','020Z')
AND bib_index.normal_heading != 'OCOLC'
AND UPPER(bib_index.display_heading) NOT LIKE '%%SET%%'
AND UPPER(bib_index.display_heading) NOT LIKE '%%SER%%'
AND bib_text.bib_id = bib_master.bib_id
AND bib_index.normal_heading IN (
    SELECT bib_index.normal_heading
    FROM bib_index
    WHERE bib_index.index_code IN ('020N','020A','ISB3','020Z')
    AND UPPER(bib_index.display_heading) NOT LIKE '%%SET%%'
    AND UPPER(bib_index.display_heading) NOT LIKE '%%SER%%'
    AND bib_id IN (
        SELECT DISTINCT bib_index.bib_id
        FROM bib_index
        WHERE bib_index.index_code IN ('020N','020A','ISB3','020Z')
        AND bib_index.normal_heading IN (%s)
        AND bib_index.normal_heading != 'OCOLC'
        AND UPPER(bib_index.display_heading) NOT LIKE '%%SET%%'
        AND UPPER(bib_index.display_heading) NOT LIKE '%%SER%%'
    )
)
ORDER BY bib_index.bib_id

```


switch to notebook

RDBMS DDL - schema, E-R models

RDBMS schema / models

- ❖ how do we decide which models go where?
- ❖ when are sets of attributes one relation, or two, or more?
- ❖ design methodologies, normalization

RDBMS schema design methods

- ❖ functional analysis of requirements
- ❖ clear determination of **cardinality** of relationships:
 - ❖ one to one, one to many, many to many
- ❖ integration with developer and analyst toolkits

RDBMS schema normalization

- ❖ "normal forms" - series of increasingly stringent schema design requirements that build on each other:
 - ❖ first normal form (1NF) - atomic values, one table per relation type, primary keys
 - ❖ second normal form (2NF) - eliminate dependencies
 - ❖ third normal form (3NF) - even fewer dependencies
 - ❖ Boyce-Codd normal form (BCNF) - fewer still

RDBMS schema normalization

- ✧ great examples in Wikipedia under:
- ✧ https://en.wikipedia.org/wiki/Database_normalization

RDBMS referential integrity

- ❖ schema-defined constraints within and between tables must be enforced
- ❖ no attribute values outside of acceptable datatype domain, range
- ❖ no foreign key references without primary key definitions
- ❖ consistent implementation of cardinality rules
- ❖ defined handling of cascading updates

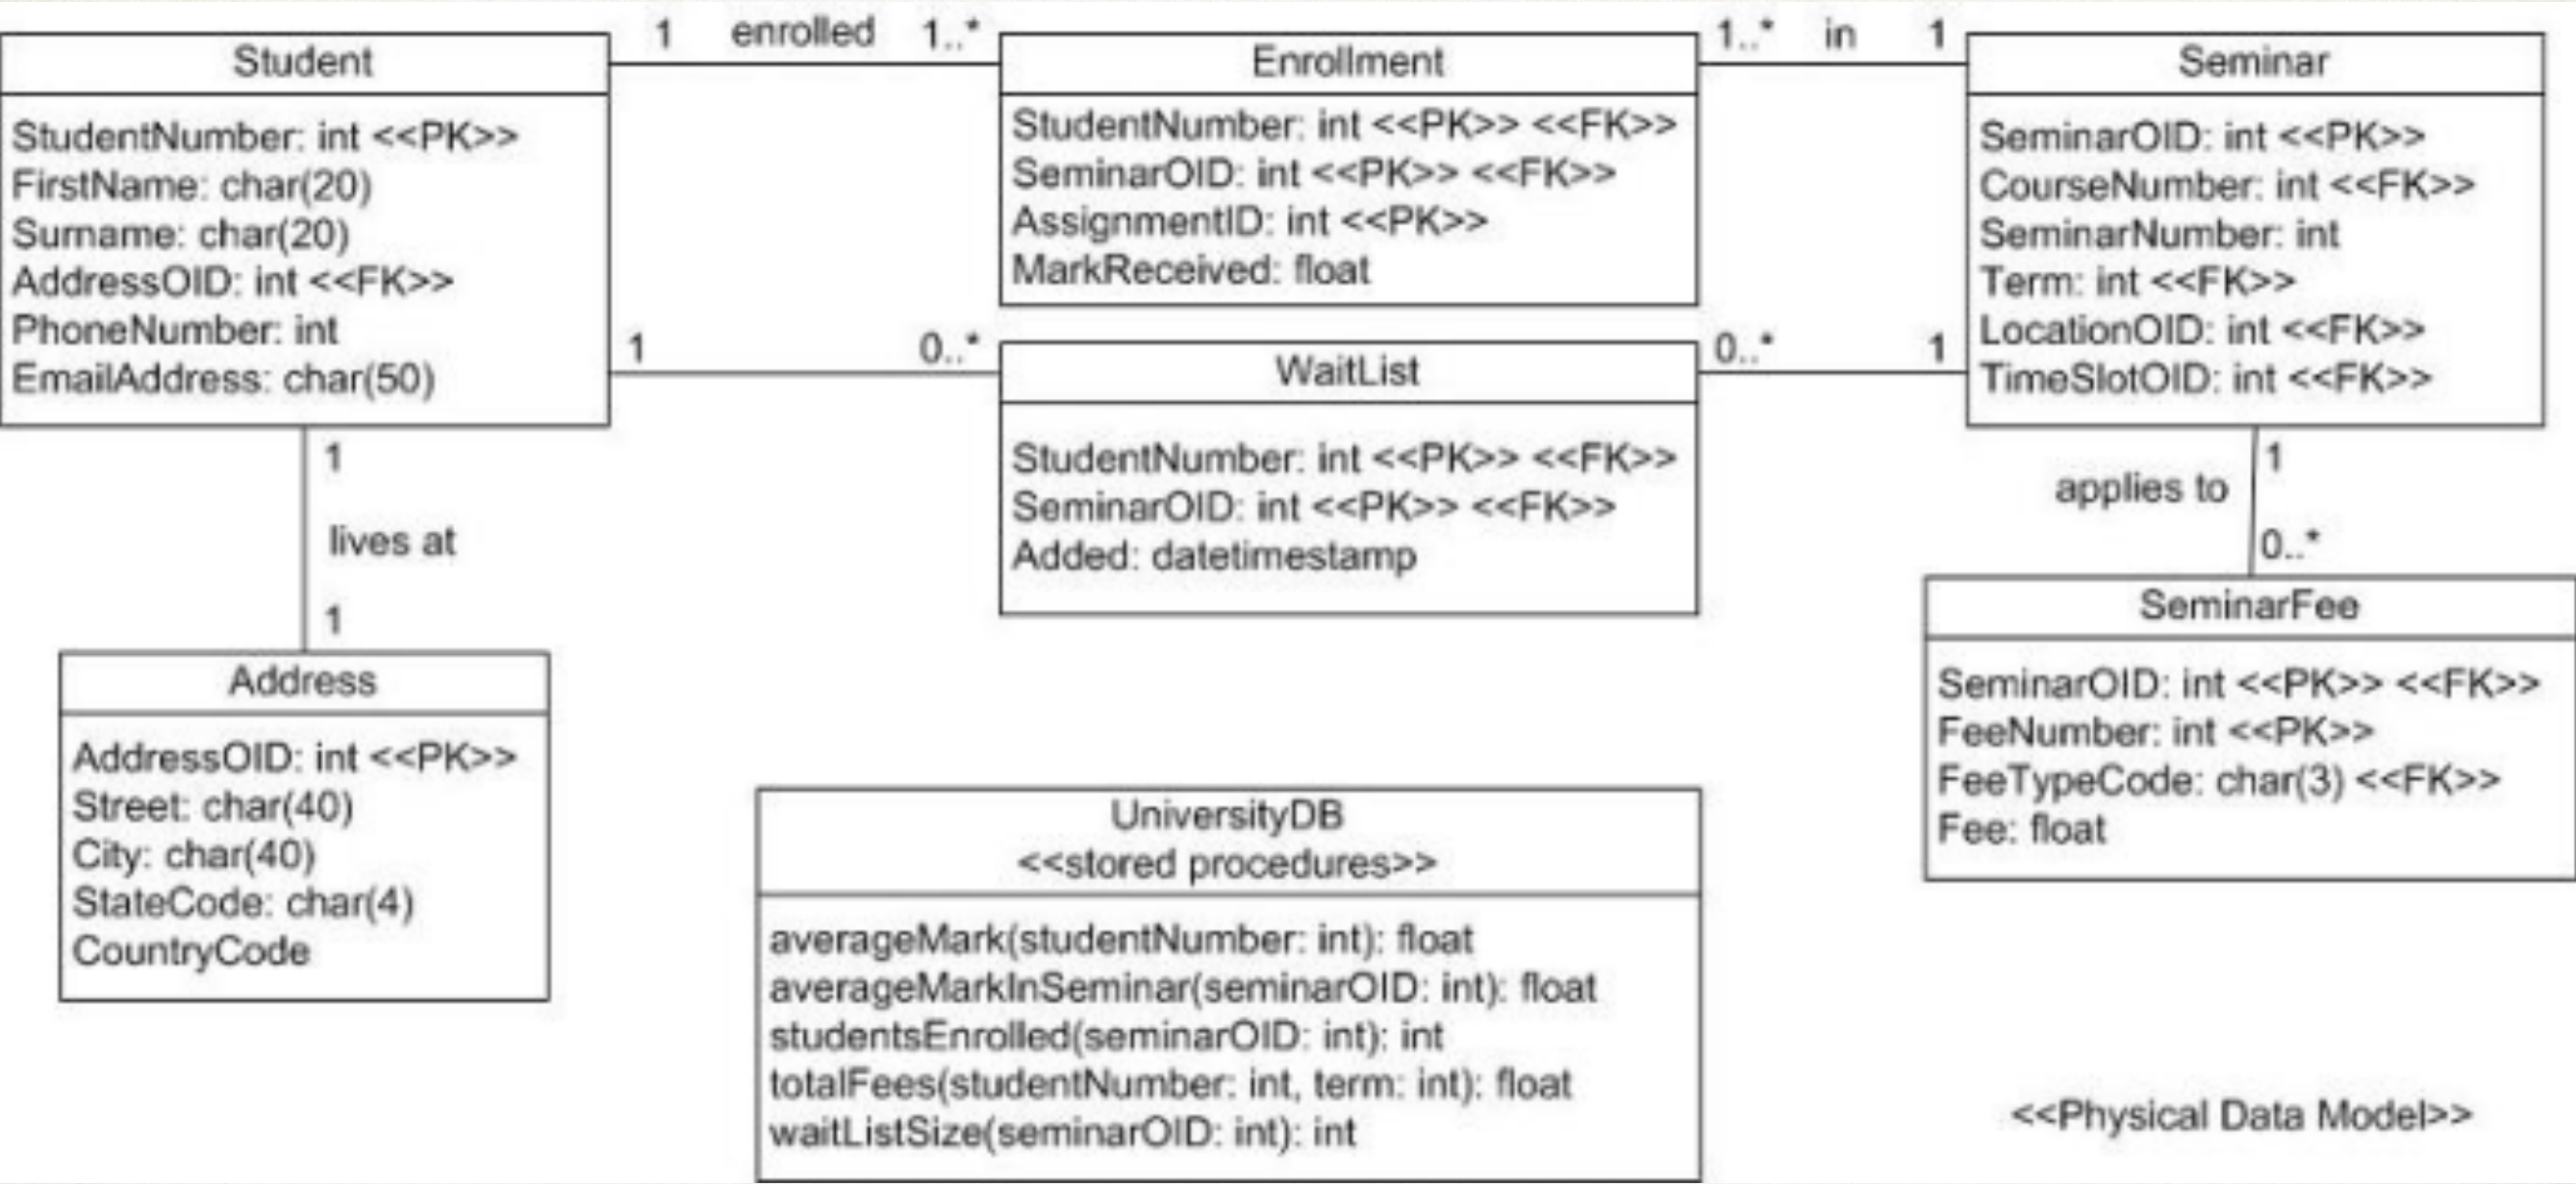
RDBMS normalization - drawbacks

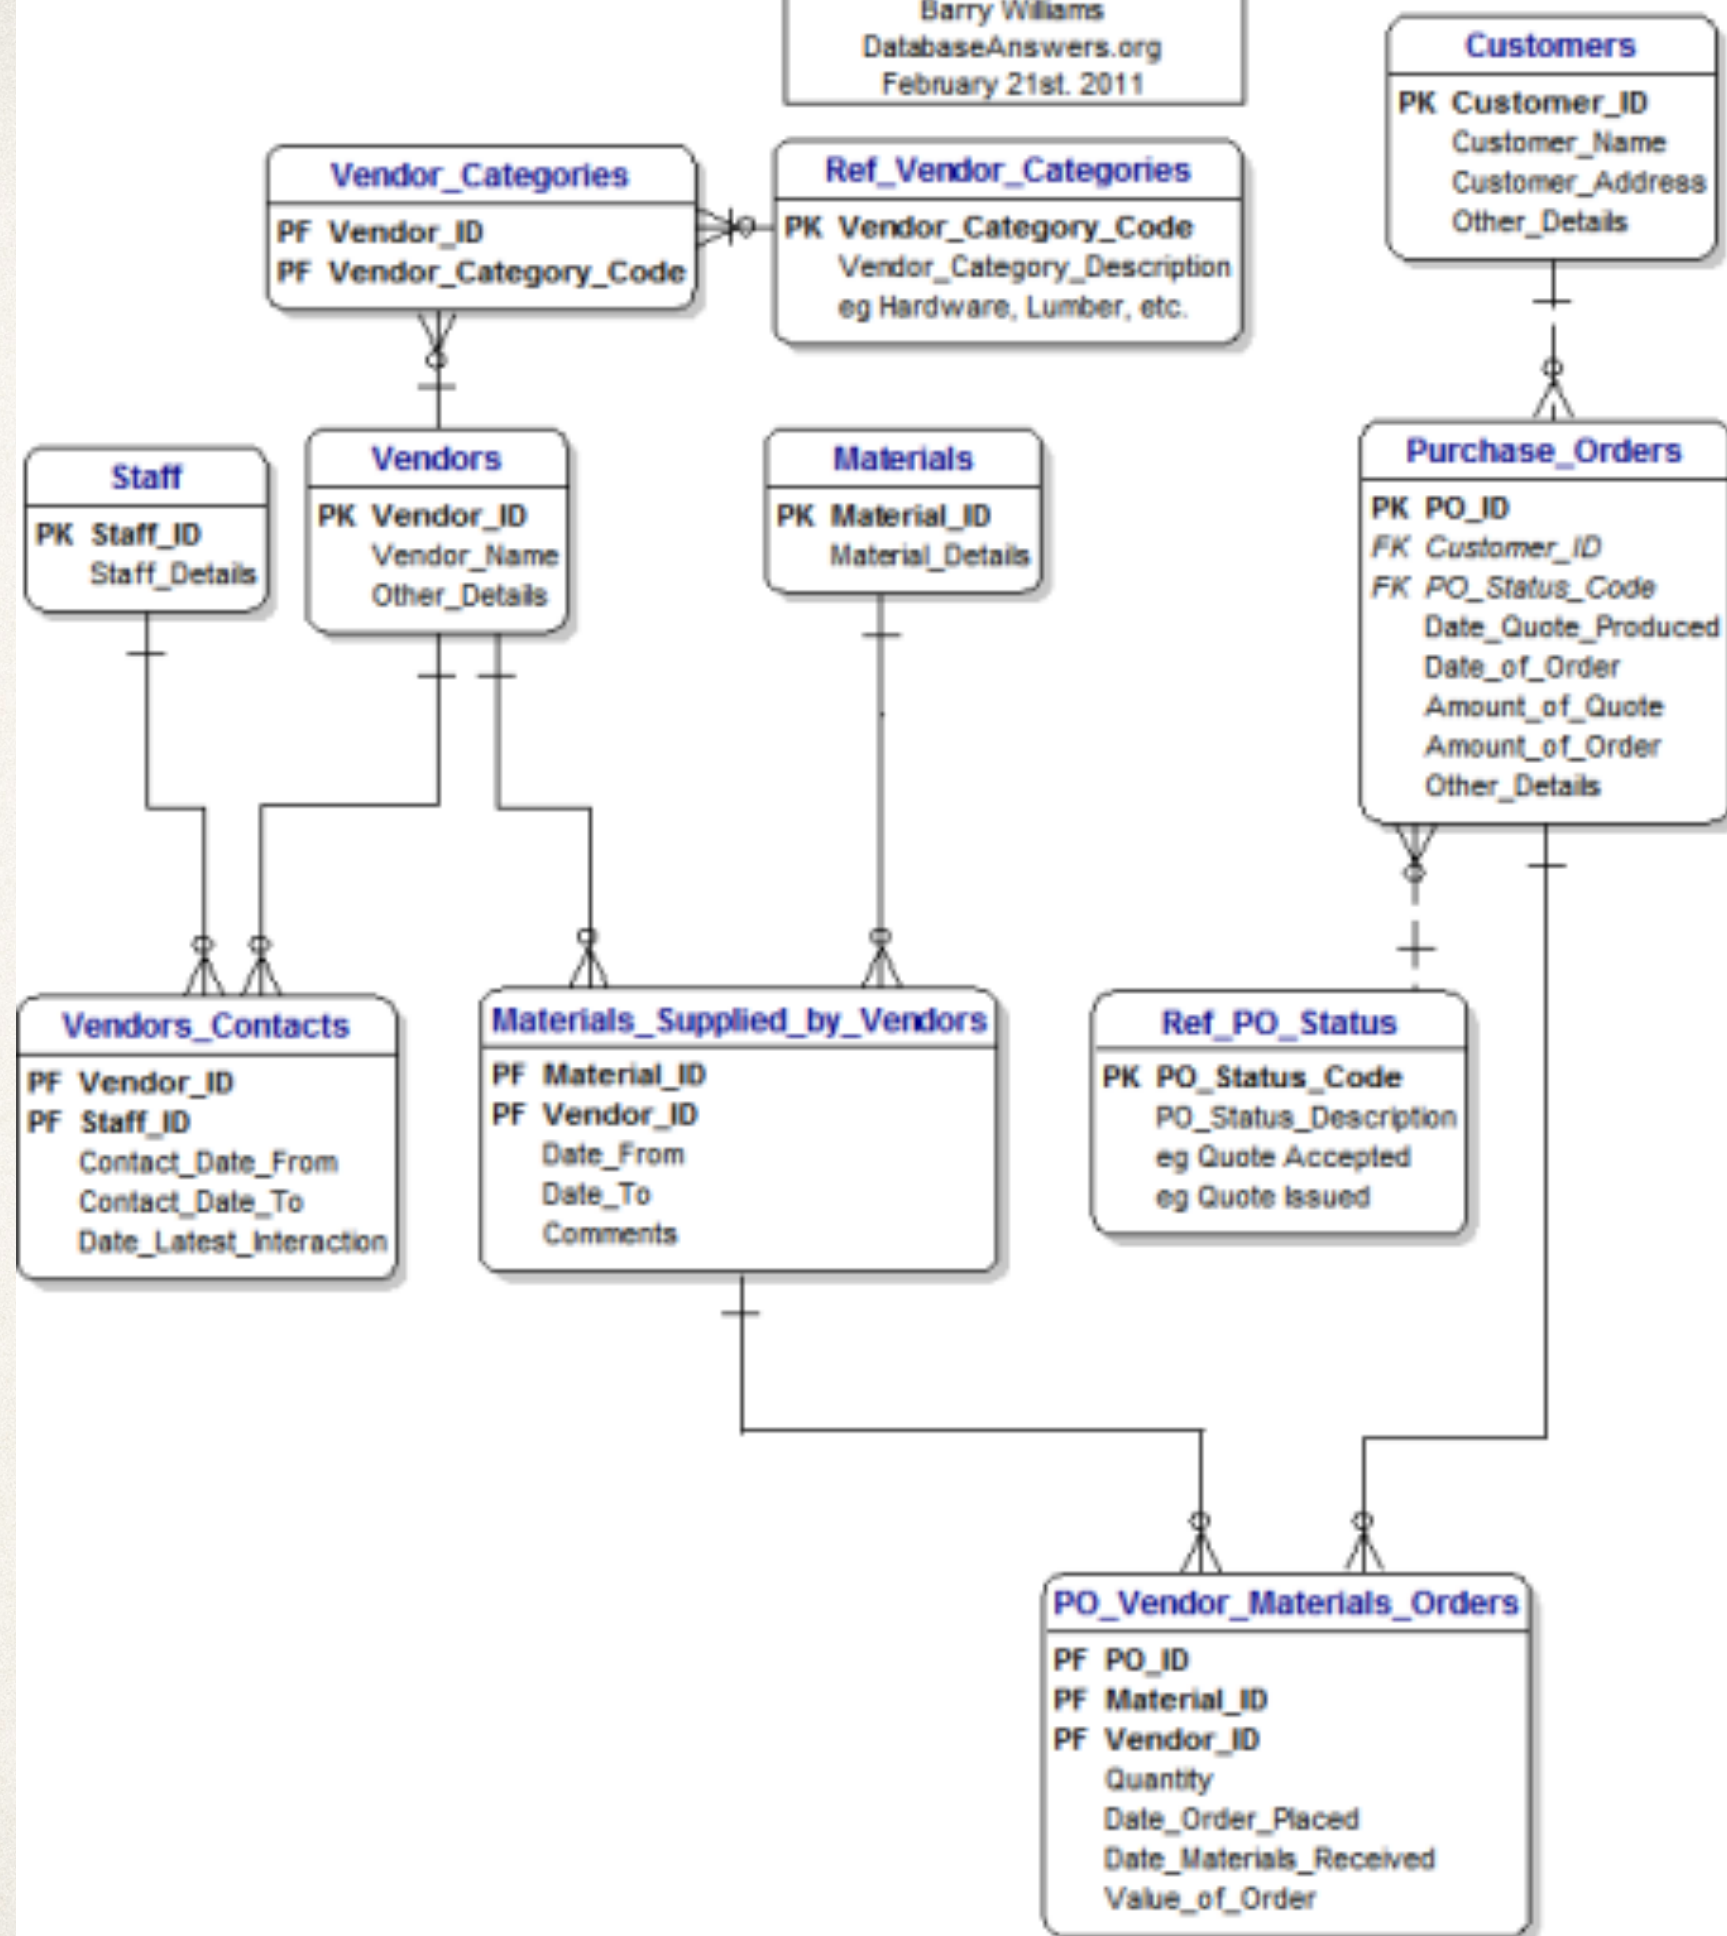
- ❖ normalization optimizes for operational applications:
 - ❖ storage, **write** (insert / update), integrity efficient
 - ❖ scales to very large transaction volumes
 - ❖ used everywhere to record business transactions
- ❖ lots of work to balance usefulness and soundness
- ❖ not always ideal for efficient analysis (**read** heavy)

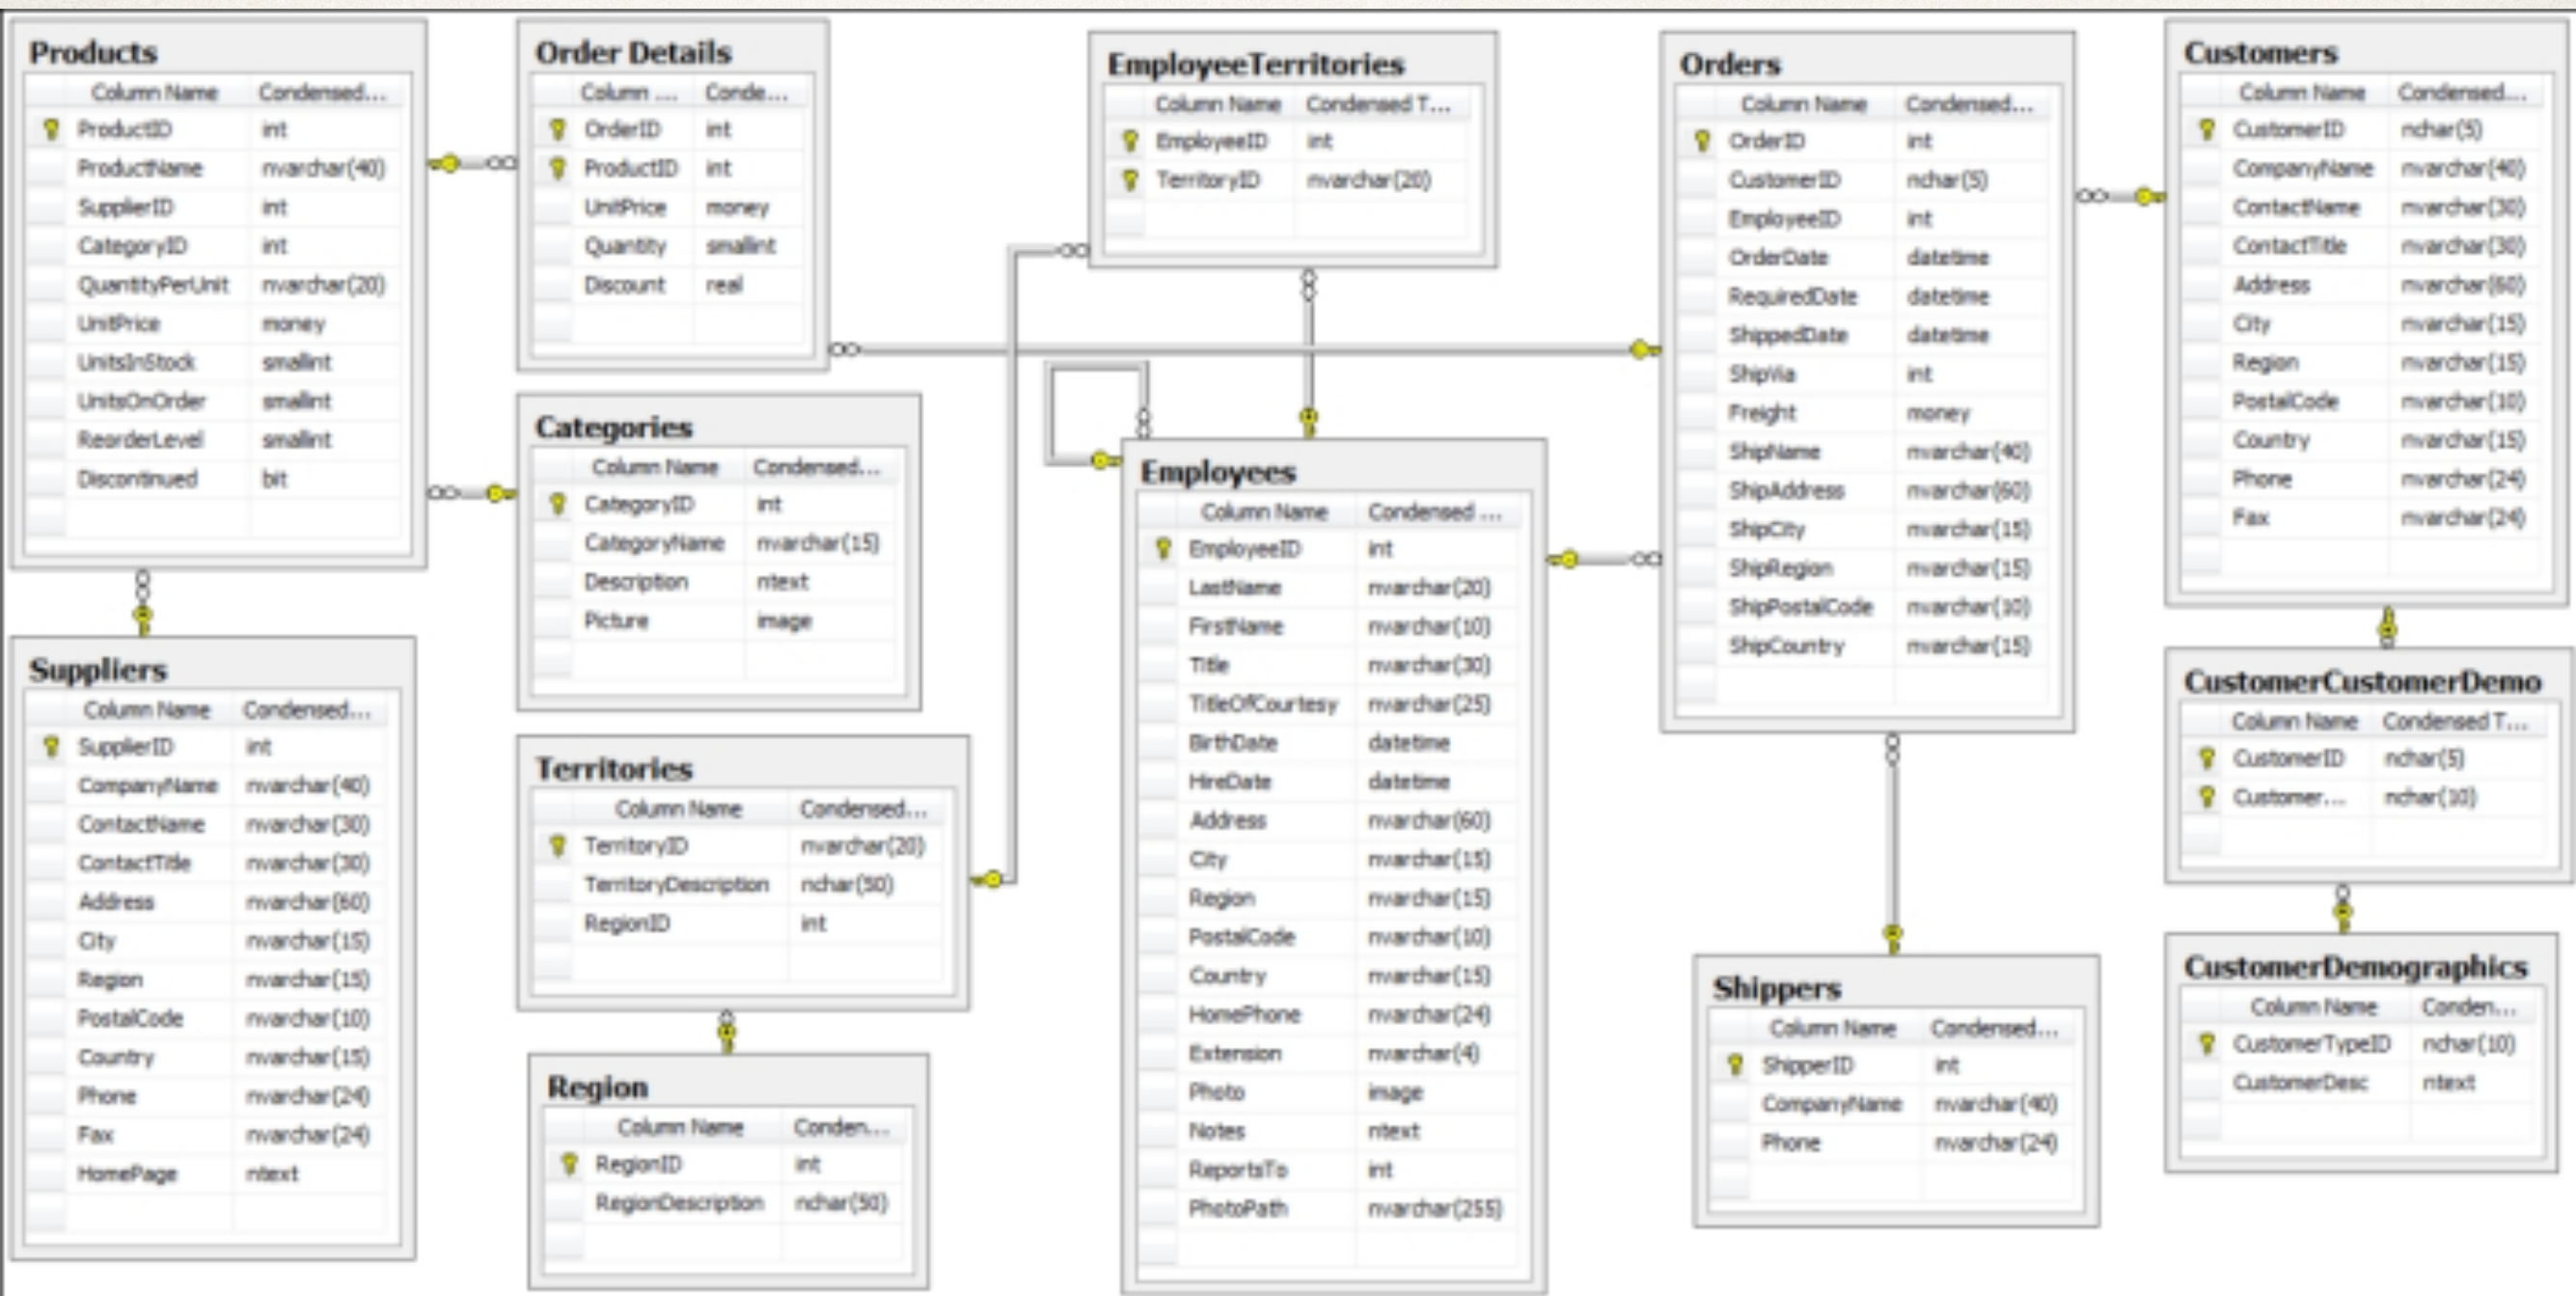
Entity-Relationship models

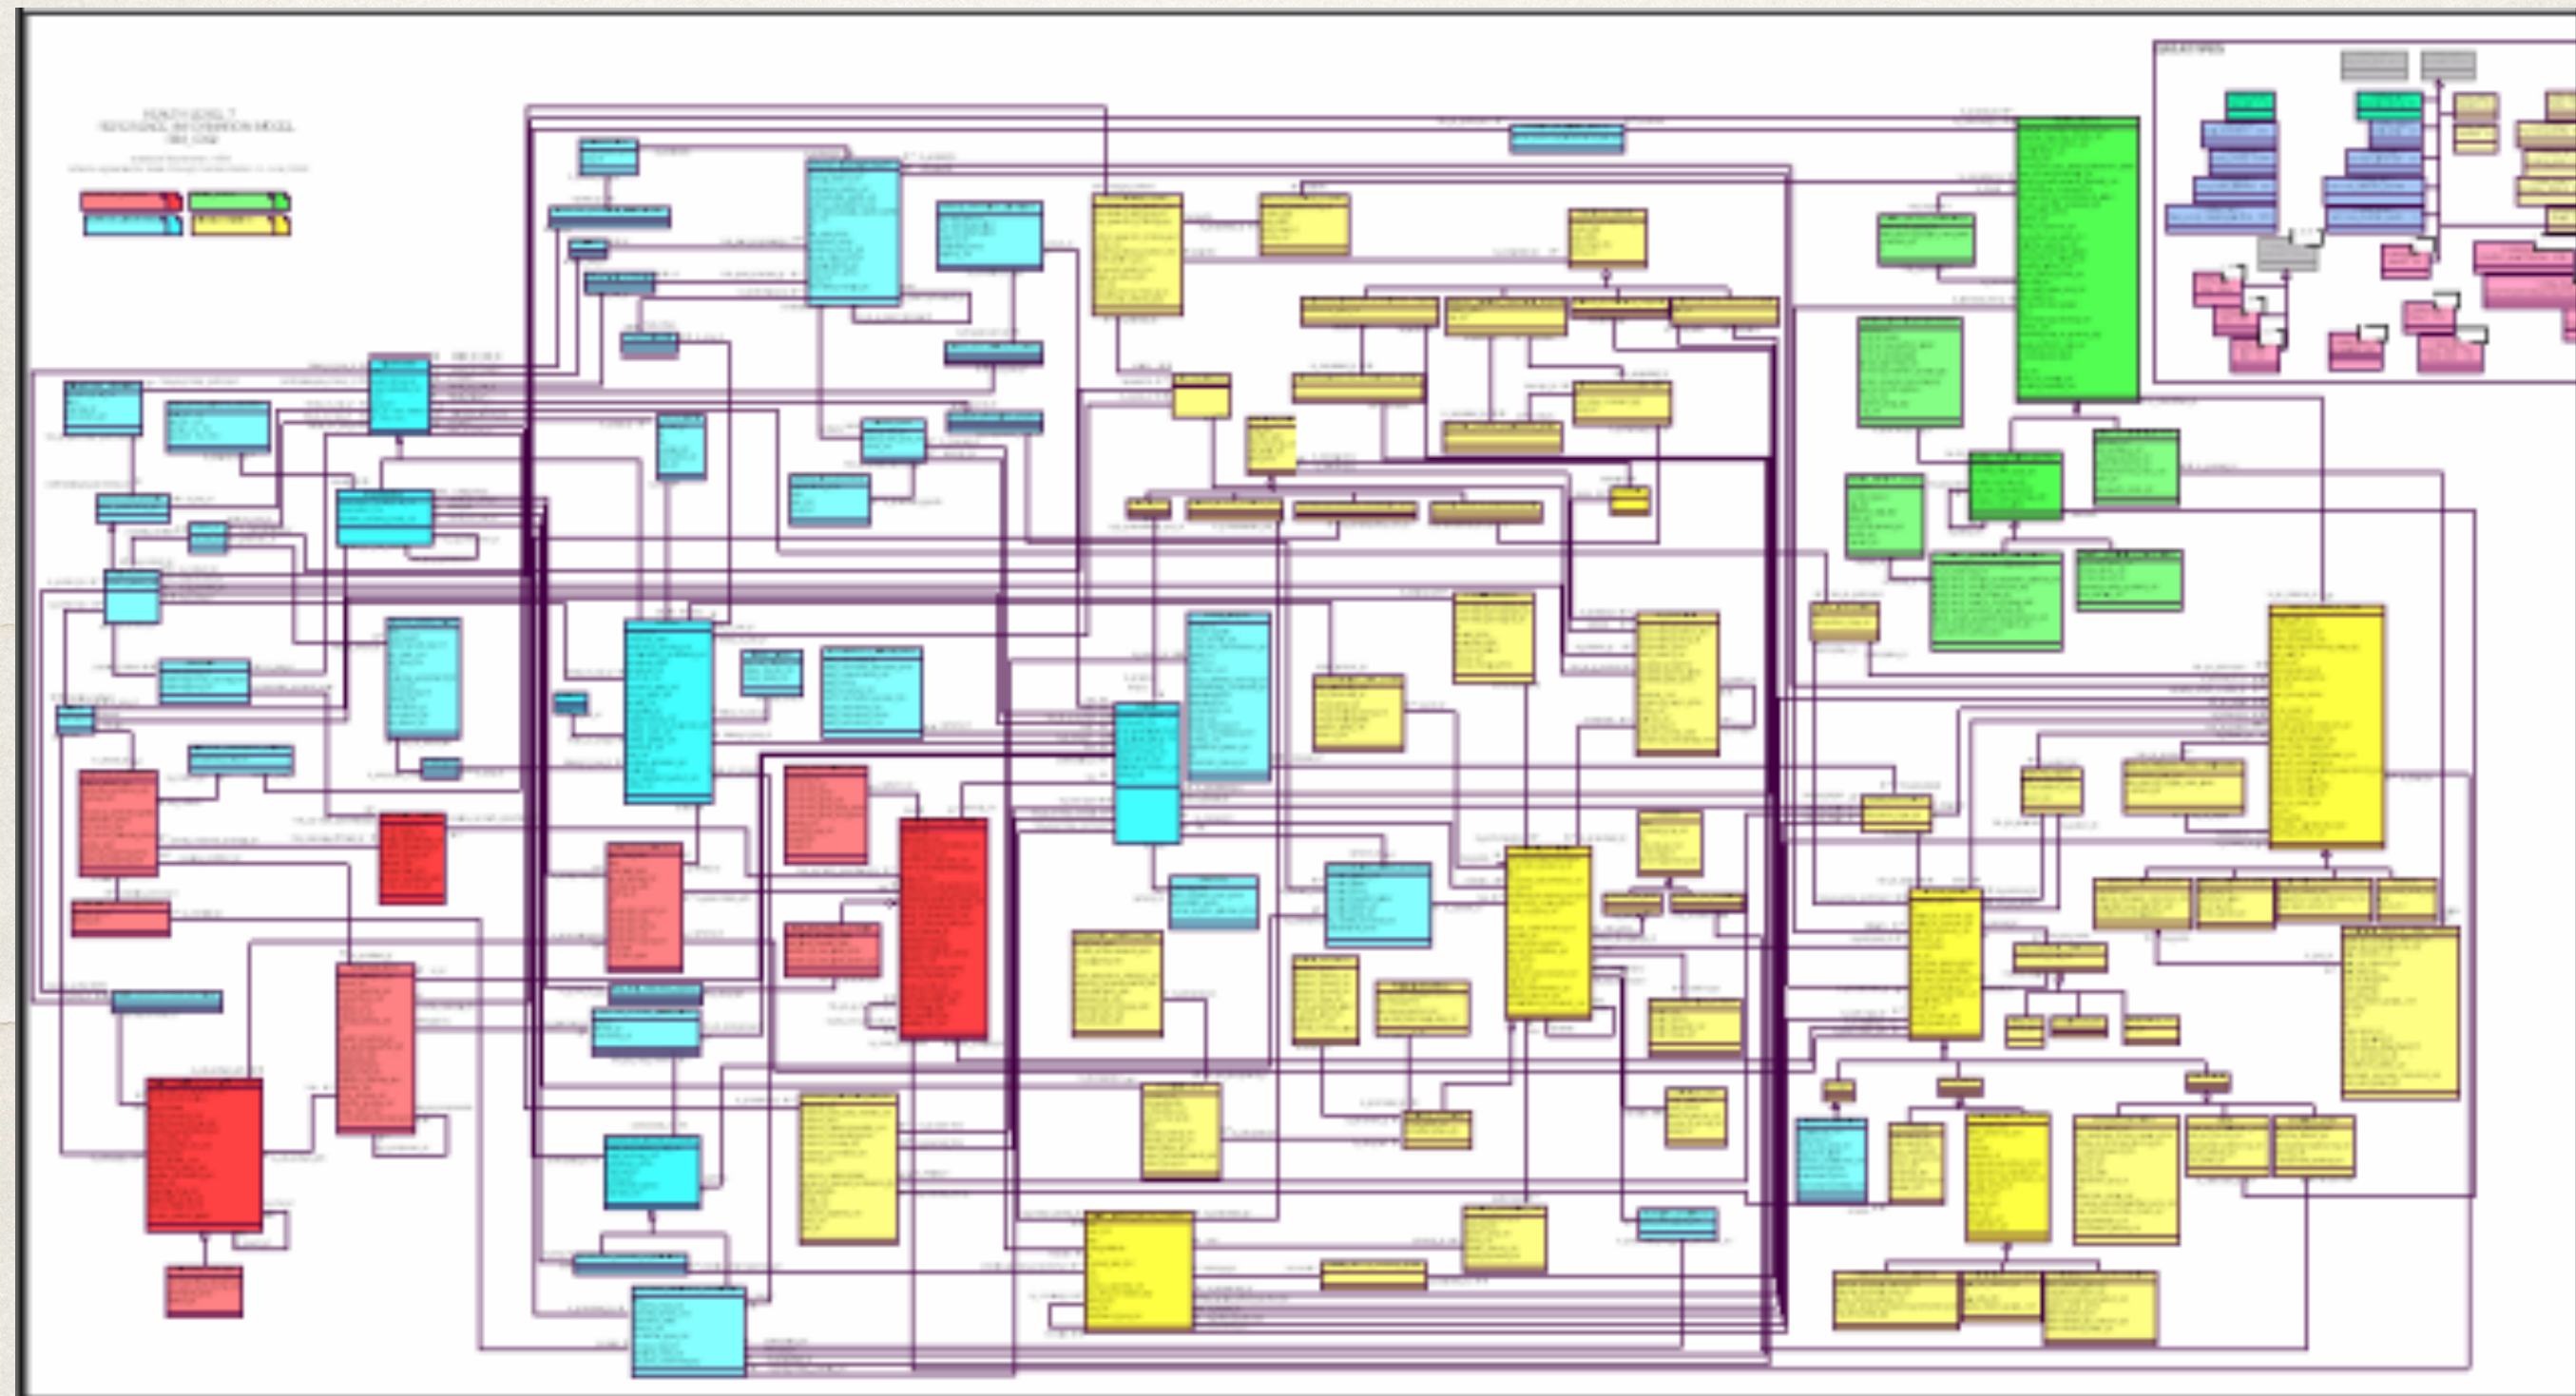
- ❖ concise visual form of representing schema
- ❖ focus on relation types (tables) and their inter-relationships
- ❖ highlight primary keys, foreign keys, cardinality
- ❖ ideally include data types, constraints if room allows
- ❖ many variations on "visual language" for these

Examples: E-R Diagrams









very common to see full designs like this

RDBMS DDL - SQL CREATE

SQL CREATE

- ❖ define new relation (table)
- ❖ simple syntax:
 - ❖ table name, attributes, types, constraints
- ❖ must be logical (order matters)

SQL CREATE examples

- ❖ CREATE TABLE Person(ident TEXT, personal TEXT, family TEXT);
- ❖ CREATE TABLE Site(name TEXT, lat REAL, long REAL);
- ❖ CREATE TABLE Visited(ident INTEGER, site TEXT, dated TEXT);
- ❖ CREATE TABLE Survey(taken INTEGER, person TEXT, quant REAL, reading REAL);

SQL CREATE - examples (2)

```
CREATE TABLE Survey(  
    taken INTEGER NOT NULL, -- where reading taken  
    person TEXT,           -- may not know who took it  
    quant REAL NOT NULL,   -- the quantity measured  
    reading REAL NOT NULL, -- the actual reading  
    PRIMARY KEY (taken, quant),  
    FOREIGN KEY (taken) REFERENCES Visited(ident),  
    FOREIGN KEY (person) REFERENCES Person(ident)  
);
```


diagram vs. schema vs. model

- ❖ E-R diagrams provide high-level overviews but might lack specifics
- ❖ DDL schema code provides low-level details but lacks high-level summary
- ❖ both are necessary in all but trivial databases

RDBMS DML - SQL INSERT, UPDATE, DELETE

SQL INSERT

- ❖ add rows to a table
- ❖ attributes must align (explicitly or implicitly)
- ❖ must abide by table definition
- ❖ may insert many rows at once

SQL INSERT examples

- ❖ CREATE TABLE Site(name TEXT, lat REAL, long REAL);
- ❖ INSERT INTO Site VALUES ('DR-1', -49.85, -128.57);
- ❖ INSERT INTO Site VALUES ('DR-3', -47.15, -126.72);
- ❖ INSERT INTO Site VALUES ('MSK-4', -48.87, -123.40);
- ❖ INSERT INTO Site (lat, long, name) VALUES (-49.85, -128.57, 'DR-1'), (-47.15, -126.72, 'DR-3'), (-48.87, -123.40, 'MSK-4');

SQL INSERT examples (2)

- ❖ CREATE TABLE JustLatLong(lat text, long text);
- ❖ INSERT INTO JustLatLong SELECT lat, long FROM Site;

SQL UPDATE

- ❖ change existing records
- ❖ won't add new or delete existing records
- ❖ must abide by schema constraints
- ❖ can use subqueries to extract or constrain values from data located elsewhere
- ❖ easy to make mistakes!

SQL UPDATE - examples

❖ compare:

❖ UPDATE Site SET lat = -48.87, long = -125.40;

❖ UPDATE Site SET lat = -48.87, long = -125.40
WHERE name = 'MSK-4';

❖ what effects can forgetting a constraint have?

SQL DELETE

- ❖ remove existing records
- ❖ must abide by integrity constraints
- ❖ can use subqueries to extract or constrain values from data located elsewhere
- ❖ just as easy to make mistakes!

SQL DELETE - examples

- ❖ compare:
 - ❖ DELETE FROM Site;
 - ❖ DELETE FROM Site WHERE name = 'MSK-4';
- ❖ what effects can forgetting a constraint have?

switch to notebook

RDBMS in practice: transactions, functions, triggers

Transactions

- ❖ "ACID" properties:
 - ❖ Atomicity - all steps complete or all fail
 - ❖ Consistency - db is consistent whether transaction succeeds or fails
 - ❖ Isolation - two simultaneous transactions follow rules for each to complete independently
 - ❖ Durability - survive system failure, etc.

Transactions - example

START TRANSACTION;

UPDATE account SET balance=balance-900

WHERE account_num=9001;

UPDATE account SET balance=balance+900

WHERE account_num=9002;

COMMIT;

ROLLBACK;

Functions

- ❖ allows users to define commonly used operations as functions for use in SQL statements
- ❖ similar to function definition in procedural programming
- ❖ defines name, parameters, return type
- ❖ often requires additional permissions

Functions - example

```
CREATE FUNCTION CtoF(Celsius FLOAT)
  RETURNS FLOAT
  RETURN (Celsius * 1.8) + 32;
```

```
SELECT name, CtoF(boiling_point)
FROM elements;
```


Triggers

- ❖ based on some conditions, perform specific operations
- ❖ INSERTs, UPDATEs, and DELETEs are typical trigger actions
- ❖ often used for formal logging of activities and data changes

Triggers - example

```
CREATE TRIGGER Books_Delete
AFTER DELETE ON Books
REFERENCING OLD ROW AS Old
FOR EACH ROW
INSERT INTO Books_Deleted_Log
VALUES (Old.title);
```


Exercise 03
